


Introduction to Parallel Programming with Share Memory Open Multi-Processing (OpenMP)

Henry R. Moncada



August 20, 2020

CONTENIDO

INTENDED LEARNING OUTCOMES

BACKGROUND

PROGRAM 1

END

References

INTENDED LEARNING OUTCOMES

- ▶ Have basic understanding of Parallel programming with OpenMP
- ▶ Understand OpenMP core syntax: directives, constructs, parallel region
- ▶ Compile and run your first OpenMP code on your laptop or a supercomputer




Speed

BACKGROUND


TOTALLY UNRELATED CONCEPTS

OPENMP DEFINED

OpenMP is a Parallel Programming Model for Shared memory and distributed shared memory multiprocessors. The underlying architecture can be shared memory UMA or NUMA.



Uniform Memory Access (UMA)



Non-Uniform Memory Access (NUMA)

- ▶ OpenMP stand for Open Multi-Processing
- ▶ OpenMP is not a computer language.
- ▶ OpenMP is an Application Programming Interface (API) that supports multi-platform shared memory multiprocessing programming.
- ▶ Works in conjunction with C/C++ or Fortran
- ▶ Applications built using hybrid model of parallel programming:
 - ▶ Runs on a computer cluster using both OpenMP and Message Passing Interface (MPI)
 - ▶ OR through the use of OpenMP extensions for non-shared memory systems.

COMPILER DIRECTIVES AND OPENMP DIRECTIVES

A directive is a language construct that instructs the compiler on how to process its input

- ▶ Compiler directives

- ▶ Example C/C++ compiler

- ▶ `#define` substitutes a preprocessor macro
 - ▶ `#include` inserts the content of a particular file
 - ▶ `#pragma` issues special commands to the compiler

- ▶ Examples in Fortran

- ▶ `!DIR$` issues special commands to the compiler

- ▶ OpenMP Directives are used to express parallelism in OpenMP

- ▶ In Fortran directives begin with

```
!$omp  
c$omp  
*$omp
```

- ▶ In C/C++ directives begin with

```
#pragma omp
```

```
#pragma omp parallel
```

The `omp` keyword signals the pragma as OpenMP specific. Non OpenMP compilers will ignore.

```
!$omp parallel
```

In fixed form, a line beginning with one of the above keywords and containing a space or zero in the sixth column will be treated as an OpenMP directive. It will be treated as a comment by non-OpenMP compilers.

OPENMP CONSTRUCT, CLAUSES AND PARALLEL REGION

- ▶ A **construct** is a specific OpenMP executable directive
- ▶ OpenMP directives may include various **clauses** to provide further information on the expected behaviour of the OpenMP implementation

```
!$omp construct [clause [clause ]...]
```

Example Core Syntax in Fortran

!\$omp parallel num_threads(4)

Compiler directive Construct Clause

- ▶ A **parallel region** is a region executed by all threads
 - ▶ Default storage attributes are defined by a data environment (we come back to this later)
- ▶ In C/C++ a **parallel region** is included between `{ ... }` after a directive
- ▶ In Fortran a **parallel region** is included between a directive pair

```
#pragma omp parallel  
{  
  
...  
  
}
```

```
!$omp parallel  
  
...  
  
!$omp end parallel
```

OPENMP PROGRAMS REQUIRE INCLUDING A HEADER FILE

OpenMP programs require including a header file:

Language	Header Files
Fortran 77	INCLUDE 'omp_lib.h'
Fortran 90	use omp_lib
Fortran 95	use omp_lib
C	#include <omp.h>
C++	#include <omp.h>

QUESTION

Which thread is going to execute the subroutine hello() ?

```
program hello.world
  use omp_lib
  implicit none

  !$omp parallel
    call hello(omp_get_thread_num()) ! obtain thread number
  !$omp end parallel

end program
```

```
subroutine hello(tid)
  integer tid
  write(*,*) 'hello world !!!, from thread = ', tid
end subroutine
```

```
#include <stdio.h>
#include <omp.h>

int hello(int);

int main() {
  #pragma omp parallel
  {
    hello(omp_get_thread_num()); // obtain thread
                                number
  }
  return 0;
}
```

```
int hello(int tid){
  printf( "Hello, World !!!, From thread %d!\n", tid);
}
```

QUESTION

Which thread is going to execute the subroutine hello() ?

```
program hello.world
  use omp_lib
  implicit none

  !$omp parallel
    call hello(omp_get_thread_num()) ! obtain thread number
  !$omp end parallel

end program
```

```
subroutine hello(tid)
  integer tid
  write(*,*) 'hello world !!!, from thread = ', tid
end subroutine
```

```
#include <stdio.h>
#include <omp.h>

int hello(int);

int main() {
  #pragma omp parallel
  {
    hello(omp_get_thread_num()); // obtain thread
                                number
  }
  return 0;
}
```

```
int hello(int tid){
  printf( "Hello, World !!!, From thread %d!\n", tid);
}
```

Answer : The directives

- ▶ Fortran: !\$omp parallel and !\$omp end parallel
- ▶ C: #pragma omp parallel { ... }

creates a section of code that is run from all available threads.

OPENMP ENVIRONMENT VARIABLES

Important environment variables to be used when you compile code with OpenMP

- ▶ Set the number of threads using the environment variable `OMP_NUM_THREADS`.

- ▶ For the `csh` or `tcsh` shell, enter:

```
set OMP_NUM_THREADS= <number of threads to use>
```

- ▶ For the `bash` shell, enter:

```
export OMP_NUM_THREADS= <number of threads to use>
```

- ▶ When executing a program (`a.out`)

```
env OMP_NUM_THREADS= <number of threads to use> ./a.out
```

- ▶ Process binding is enabled if this variable is `true` ... i.e. if `true` the runtime will not move threads around between processors.

```
OMP_PROC.BIND true | false
```

- ▶ Control how “omp for schedule(RUNTIME)” loop iterations are scheduled

```
OMP_SCHEDULE schedule [chunk.size]
```

OPENMP HELLO WORLD

EXERCISE 1: OUR FIRST OPENMP PROGRAM

```
#include <stdio.h>
#include <omp.h>

int main(int argc, char *argv[]) {

#pragma omp parallel
{
    int tid = omp_get_thread_num();
    printf("Hello World from thread %d!\n", tid);
}
return 0;
}
```

OpenMP include file

Parallel region with default
number of threads

Library function to return thread ID

end of parallel region

```
program hello_world
use omp_lib

!$omp parallel
write(*,*) 'Hello World from thread', omp_get_thread_num()
!$omp end parallel

end program hello_world
```

OpenMP fortran module

Parallel region with default
number of threads

Library function to return thread ID

end of parallel region

Hello_World.c

Hello_World.f90

```
#include <omp.h>
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {
    int nthreads, thread_id;

    /* Fork a team of threads giving them their own copies of
       variables */
    #pragma omp parallel private(nthreads, thread_id)
    {
        /* Obtain thread number */
        thread_id = omp_get_thread_num();
        printf("Hello World from thread = %d\n", thread_id);

        /* Only master thread does this */
        if (thread_id == 0) {
            nthreads = omp_get_num_threads();
            printf("Number of threads = %d\n", nthreads);
        }
        /* All threads join master thread and disband */
    }
    return 0;
}
```

with GNU/GCC compiler

```
$ gfortran -fopenmp -o out omp_hello.f90
```

```
$ export OMP_NUM_THREADS=16
$ ./out
```

with GNU/GCC compiler

```
$ gcc -fopenmp -o out omp_hello.c
```

with LLVM/CLANG compiler

```
$ clang -fopenmp -o out omp_hello.c
```


END

QUESTIONS

REFERENCES

► OpenMP specifications:

<http://www.openmp.org/specifications/>

► OpenMP summary Card:

► C:

<http://www.openmp.org/wp-content/uploads/OpenMP-4.0-C.pdf>

► Fortran:

<http://www.openmp.org/wp-content/uploads/OpenMP-4.0-Fortran.pdf>